

SSM 2026 Talk — Working Session Transcript

A reconstructed transcript of the collaborative session between **Derek C. Briggs** and **Claude** (Claude Code, Opus 4.8) that produced the SSM 2026 talk materials: the *verification-frontier* simulation (`verification_frontier.Rmd / .py`) and a working slide deck. Later exchanges refine the model (a detection-vs-appraisal reframe, the verification-cost axis, an item-characteristic-curve explainer, a Panel B fix) and publish the auditable simulation to GitHub.

Talk: "Is Measurement Necessary for AGI — or Are We Just Making a Bad Joke?" Society for the Study of Measurement, Edinburgh, 24 June 2026.

Derek (user)

Read CLAUDE.md and act upon instructions.

[Project brief (CLAUDE.md): help Derek Briggs brainstorm and focus a 20-minute talk for the Society for the Study of Measurement (SSM), Edinburgh, 24 June. Inputs provided: (1) the January 2026 SSM proposal; (2) a ChatGPT 5.2 brainstorm transcript on measurement/RLHF/LLM helpfulness; (3) background on how Derek conceives measurement -- the 2021 book and the Educational Measurement 5th ed. Ch. 21 (TISM); (4) how AI researchers think about measuring intelligence -- Chollet, "On the Measure of Intelligence," and "The Leaderboard Illusion" (Chatbot Arena). Constraint: ~3 days, conceptual + possible simulation, no empirical study.]

Claude (assistant)

[Read the proposal, the brainstorm transcript, and extracted the TISM chapter, Chollet, and the Leaderboard Illusion.]

I've now worked through all six inputs. Before we start honing, let me reflect back the live intellectual material and then ask a few questions.

What I think you actually have

Your proposal contains **two distinct theses**, and I think only one is the keeper for a conceptual talk.

Thesis 1 (the TISM audit): apply T-I-S-M to the SFT/RLHF pipeline and show it fails on all four counts -- no substantive theory of the construct, instrumentation that echoes Thurstone without his care, no scale/unit/reference standard, no uncertainty quantification beyond temperature. Correct, well-supported, and the **easy** part. The Leaderboard Illusion hands you the empirical ammunition (Goodhart, no anchoring, silent deprecation breaking Bradley-Terry, drift). Chollet hands you the conceptual ammunition (skill != intelligence maps onto your **testing != measurement**).

Thesis 2 (the inversion): the original one, buried in your last paragraph. Measurement exists to characterize variability; generative AI development is the first large-scale enterprise whose explicit goal is to **eliminate** variability. So the question isn't merely misuse of a word -- it's whether measurement can be the right frame for an artifact engineered to erase its own variance.

The ChatGPT transcript is a third strand: you worked out a sound IRT/BTL design (models-as-persons, prompts-as-items, raters-as-facets) and then hit exactly the wall that proves Thesis 2.

My recommendation

Make **Thesis 2 the spine**, use the **TISM audit as scaffolding**, and bring in Chollet + the Arena paper as your two exhibits. Payoff line:

AGI development inverts the precondition of measurement. We may need to recognize that what's happening is not measurement at all, but a control/optimization activity wearing measurement's clothes.

Where a simulation could earn its place

Simulate the inversion: in a Bradley-Terry/Elo setup, show that as model true-scores converge (variance $\rightarrow 0$), Fisher information per comparison collapses and rankings destabilize -- the scale degenerates because the engineering succeeded.

Clarifying questions

1. **Center of gravity** -- critique, constructive offering, or *mirror* (AI forces a new insight about measurement itself)? I advocate the third.
2. **The Claude Code angle** -- motivating frame, or set aside to stay in the RLHF space?
3. **Simulation appetite** -- yes to the convergence/degeneracy simulation, or purely conceptual?
4. **The joke** -- keep humor as the running example, or switch to helpfulness/coding?

My instinct: open with the Claude Code lived experience, run Thesis 2 as the spine, keep the joke as the recurring example, include one simulation.

Derek (user)

I kind of like the mirror perspective -- can measurement even be the right epistemic frame for an artifact engineered to erase its own variance? But let me push you and myself.

I've argued that in learning to use AI as a trusted collaborator it's important to run within-person experiments, using AI on a task where you have expertise or where you compare models / test sensitivity to context -- easiest where outcomes are verifiable. At the individual level I don't think we need measurement per se. And holding the prompt constant, the only source of variability in LLM response is temperature. But aren't there almost an infinite number of ways to prompt an AI to accomplish a complex task? Maybe measurement matters -- in measuring the ability of a person to interact with an AI effectively. So maybe the point is: speaking of measuring latent attributes of AI models is a bad idea; instead, our focus should be on measuring AI Literacy constructs for the people using the models?

Claude (assistant)

This is the right kind of pushback, and I think you've just found the actual talk.

Name the move

You're relocating the variability to where it actually lives. That *rescues* measurement rather than dissolving it -- and completes the mirror: take the model seriously as a measured object, discover its variance has been engineered away, which forces you to ask where the variance went. Answer: into the human. It also overturns the default in your own transcript (GPT ranked "measuring the user" *least* defensible).

A correction that's a gift

Your Ch. 21 traces measurement's two historical roots: **intra**-individual (Fechner/Thurstone) and **inter**-individual (Galton/Spearman). The LLM-as-object is degenerate on **both**: intra-variance is just temperature (a dial you set to zero); inter-variance is ~5 converging frontier models. So the model is not a measured object -- it's **instrumentation**. TISM role reassignment: object = the human; instrument = model+config; item = prompt+verifiable goal; score = verifiable task success.

Three pushes

1. **Does "AI Literacy" have a T?** It's a nascent, in-flux construct -- the hardest category in Ch. 21. Owe a substantive theory before measuring.
2. **The Chollet confound.** Measure skill-acquisition efficiency *net of priors and experience*. "Good with AI" is confounded with "knows the domain." Tool fluency vs. domain priors -- a DIF/facet problem.
3. **The crown jewel:** you measure the human *through* an instrument deliberately engineered to minimize the construct variance you want to detect. A good model makes a novice's output look expert. Penmanship in the age of the word processor; navigation in the age of GPS.

The forward hook

As models improve, does AI use **compress** or **amplify** inter-person variance on verifiable tasks? That's the empirical jump-off.

Proposed spine

Open with Claude Code -> model fails on both pillars (mirror) -> model is instrument, not object -> relocate to critical AI literacy -> the instrument compresses the construct -> close on compression-vs-amplification.

Two questions: (1) AI Literacy general or domain-bounded? (2) How much weight on the critique foot? I lean fast critique, long reframe.

Derek (user)

(1) I like the spine. Let me point you to where I've defined AI Literacy: a 4-part blog series. If there's a construct it hinges on the ability to use AI *critically*. The general "habits of mind" are domain general, but developed domain-specifically. Your "forward hook" is not so much about dimensions of an AI Literacy construct but about expertise and skill development if everyone is already AI literate at a base level -- if all have access to AI, it becomes hard to discern individual differences. (2) Yes, fast critique, long reframe. (3) Open to a simulation -- that's what you're good at. Also: joke as the star, one quick verifiable foil.

Claude (assistant)

[Built simulations/verification_frontier.py and rendered a 3-panel figure.]

The model: one task = a Bernoulli outcome, $P(\text{success} \mid \theta) = c + (1 - c) * g(\theta)$, with $g = \text{logistic}(a(\theta - \delta(s)))$. θ = critical literacy; c = model capability; s = seductiveness -> catch-difficulty $\delta(s)$. Per-task value = $E_{\theta}[\text{Fisher information}]$.

- **Panel A** -- the verification frontier (ridge): information about critical literacy peaks at intermediate seductiveness, dies where the model is strong and errors obvious, and where errors are too subtle to

catch.

- **Panel B** -- the window narrows as models improve.
- **Panel C** -- per-task information collapses along the model-progress path.

Headline numbers (reliability 0.80): verifiable foil ~335 trap-items; the joke ~50; along model progress 30 -> 2,609. Two claims: critical literacy is **more** measurable where you cannot verify; and it's measurable only on a **moving frontier**.

One honest modeling choice for your judgment: c means real correctness in verifiable domains but only fluency in unverifiable ones. Want me to split the model that way? Refinement levers: correlate s with c; add a domain-prior facet; swap the reliability target.

Derek (user)

Can you generate the code as an R Markdown document so I can follow step by step? I need to make sure I'm properly modeling critical literacy in this interaction and am not just getting fooled by a high s.

Claude (assistant)

Exactly the right move -- auditing a fluent artifact you can't verify at a glance is the construct in action. I wrote simulations/verification_frontier.Rmd to be **interrogable**: every assumption flagged [CHOICE], the algebra checked numerically against finite differences, and the Fisher-information formula verified by Monte Carlo.

[Verified in R 4.5.2:]

- Check 1 (derivative): PASS.
- Check 2 (Fisher vs Monte Carlo): analytic 0.1628/0.0340/0.1041 vs MC 0.1627/0.0341/0.1042.
- Check 3 (collapse as c->1): 0.202 -> 0.096 -> 0.021 -> 0.002 -> 0.0002.
- Headline reproduced exactly (foil 335, joke 50, generations 30 -> 2,609).

Knit it and the self-checks run inline, so a changed assumption that breaks my algebra fails loudly rather than yielding a pretty wrong surface.

Derek (user)

[Screenshot of an RStudio render error.] The Rmd file didn't knit.

Claude (assistant)

Good catch by the render. crimson is a valid matplotlib color but **not** in base R's colors() table -- that's the failure. My earlier self-check ran only the computational core, not the plotting, so it missed this; a real verification gap on my part.

Fix: col = "crimson" -> col = "#DC143C" in Panel C. I then verified the full render path by tangling the Rmd and sourcing it to a graphics device (all chunks render OK, self-checks still pass). I've made full-render verification a standing practice for anything I hand you.

Derek (user)

Yes, let's do 1 then 2.

Claude (assistant)

[Implemented the verifiable/unverifiable split in the Rmd and Python; regenerated the figure.]

The cleanest mechanism keeps the verified Fisher machinery as the **map** and makes verifiable vs. unverifiable two **roads** across it, forked by one lever -- whether capability has a ground-truth signal to climb toward. Verifiable domains let $c \rightarrow \sim 1$ (collapse); unverifiable cap c well below 1 (the construct persists).

A bug I had to surface and own. My `sensitivity()` function used `<<-` and never restored the global parameters, so a verification probe running after that chunk reported numbers from polluted globals (delta -2.2 instead of -3.3). The clean, correct numbers are:

Point	c	s	info/item	items for rho=.80
Shared early model	0.50	0.45	0.131	30
Verifiable frontier	0.96	0.62	0.007	~719
Unverifiable frontier (the joke)	0.55	0.62	0.092	~50

I fixed `sensitivity()` to restore state via `on.exit`. Both domains start equally measurable; model progress forks them -- verifiable collapses, the joke persists.

The 20-minute structure (7 beats)

1. Hook -- how I actually use AI now (Claude Code).
2. Take the model seriously as a measured object -- fails on both pillars (fast critique; Chollet + Leaderboard Illusion as exhibits).
3. So the model is an instrument, not an object (TISM role table).
4. Relocate to critical AI literacy.
5. The verification frontier (simulation).
6. The sting -- the instrument compresses the construct's variance.
7. Close -- the measurement question of the AI era.

Two questions: (1) is beat 2 worth 4 minutes or tighter? (2) slide tooling -- Quarto, Beamer, PowerPoint, or Google Slides? I lean Quarto.

Derek (user)

1. We'll need to tighten beat 2. The simulation results only make sense after explaining the model setup as done in the Rmd. I have questions about the simulation, and we need to fix things that didn't render properly -- most notably Panel A, where the text label for the unverifiable line is truncated. 2. For the slides let's do both Quarto and PowerPoint so I can decide what works best.

Claude (assistant)

Fixed the Panel A truncation in both the R (filled.contour) and Python versions. The truncation came from multi-line labels anchored at road endpoints -- the verifiable endpoint sits at $c=0.96$, against the right edge. The robust fix: drop the fragile on-plot annotations and use a clean legend instead. Confirmed the R render is clean.

Plan adjustments noted: tighten beat 2; add a "model setup" sub-beat before the results (5a setup -> 5b results); slides in both Quarto and PowerPoint once the sim is locked.

Your turn -- what are your questions about the simulation?

Derek (user)

The model is basically a 3PL IRT model with fixed discrimination, or a Rasch model that also includes a lower-asymptote item parameter. I want to make sure I'm interpreting the relationship between a traditional item difficulty parameter and the difficulty as a function of "seductiveness" s . Right now s is purely a function of model generation, right? It doesn't vary as a function of the difficulty of the task or prompt?

Claude (assistant)

Your read is exactly right.

The structural identity

It is a 3PL with fixed discrimination: lower asymptote = c (capability); difficulty $b = \Delta(s)$; discrimination = a (fixed). **But the lower asymptote means the opposite of usual** -- in standard 3PL it's construct-irrelevant guessing; here it's the substantive star (model quality). Borrowed form, re-purposed parameter -- the Ch. 21 move, done in the open.

What delta's "difficulty" actually is

Δ is the difficulty of the *human's detection task*, not the AI's task. On the same logit scale as θ : $\Delta(s)$ = the critical-literacy level at which you have a 50% chance of catching an error of seductiveness s . Seductiveness measured in units of critical literacy required to catch it.

You're right

Yes -- s (hence Δ) is a function of model generation only. The model has **no item-difficulty facet**; each (c, s) is a single representative item, not a bank. A defensible design separates three quantities: AI-task difficulty -> item-specific c_i ; detection difficulty/seductiveness -> Δ_i ; and θ . Generation shifts the *distributions*; test information integrates over the bank.

One consequence worth a slide

A non-zero lower asymptote forfeits Rasch **specific objectivity**: critical-literacy comparisons between people are **not invariant** to which model/task you measured through.

Recommendation

Add a modest item-difficulty facet rather than leave it a one-item caricature. Build it, or keep the representative-item model and relabel it honestly?

Derek (user)

Yes, let's build the item-bank version into the Rmd, but before you do: detection difficulty is set by $\text{delta_of_s}(s) = \text{DELTA_MIN} + (\text{DELTA_MAX} - \text{DELTA_MIN}) * s$. If you build the item-bank version, that's what you'd adjust, right? I'm not sure I followed the intuition for how we set DELTA_MAX and DELTA_MIN, though I see the impact of narrowing the range in the sensitivity analysis.

Claude (assistant)

delta_of_s is exactly the hook, and how DELTA_MIN/MAX are set is the softest spot in the toy.

What DELTA_MIN/MAX mean

delta lives on the same logit scale as theta $\sim N(0,1)$ -- so it's denominated in **standard deviations of the critical-literacy population**:

s	delta(s) with [-3,3]	reading
0.0	-3	~everyone catches (trivial errors)
0.5	0	~50% catch -- maximally discriminating
1.0	+3	almost no one catches (near-uncatchable)

So they state how seductive the most/least seductive errors are, in "how much critical literacy you'd need to catch them." The honest part: there's no empirical anchor for +3 -- it stands in for calibrating trap-item difficulties against real raters.

How delta_of_s changes

I'd **promote** it: delta_of_s(s) becomes the *mean* of an item-difficulty distribution, $\text{delta_i} \sim \text{Normal}(\text{delta_of_s}(s), \text{SIGMA_DELTA})$, with items scattered around it. Once difficulty is a distribution relative to the population, the exact DELTA_MIN/MAX calibration becomes much less load-bearing.

Two choices to bless: (1) Normal scatter with SIGMA_DELTA = 1.0? (2) keep c_i fixed per generation this round?

Derek (user)

1. Yes, draw delta at random from a normal distribution.
 2. Yes, keep c_i fixed per generation for now.
-

Claude (assistant)

[Checked the direction empirically, then built the item-bank section.]

A scattered bank carries **less** average information than a perfectly-targeted single item (off-target prompts dilute it), so items-needed rises modestly -- joke 50 -> ~55, verifiable 719 -> ~769 -- and grows with SIGMA_DELTA, but the divergence is untouched. (Also: frontier seductiveness $\delta=0.72$ already sits past the info-maximizing difficulty ~-0.5.)

New section "An item-bank version" with three parts: **Same bank, two capabilities** (one shared difficulty bank scored through $c=0.55$ vs 0.96 -- identical trap-items informative through a weak model, near-useless through a strong one); **The divergence survives a real bank** (the C' figure); **Sampling variability** (median 54 items, 90% interval [52, 58] across 300 random banks).

Safeguards held: a stopifnot confirms the bank code reduces exactly to the representative-item model, and globals stay intact after the full run.

Derek (user)

In section 4 of the Rmd the notation for Information isn't rendering properly based on the line $\mathcal{I}(\theta)$. Can you fix this?

Claude (assistant)

The culprit is the AMS math fonts -- $\mathcal{I}$ (and $\mathbb{E}$) depend on MathJax extensions that aren't loading in your knit setup. Robust fix: drop those font commands and use plain I for Fisher information and E for the expectation -- which also matches the code's notation. Fixed all three occurrences (including the $\mathcal{N}$ in the item-bank section, same latent failure). They now use only core LaTeX, so they render regardless of the math setup.

Derek (user)

Yes [build the slides].

Claude (assistant)

[Quarto not on the CLI path here; authored the deck for you to render in RStudio.]

Built slides/ssm_talk.qmd -- one Quarto source that renders to **both** RevealJS (HTML) and PowerPoint, so you can compare. ~14 slides for 20 minutes with the two tweaks: beat 2 tightened to two slides; a "model setup" sub-beat added before the results. Figures regenerate from a compact copy of the verified model core, which carries its own stopifnot self-check, and the inline trap-item numbers (30 / 719 / 50) compute live. Every R chunk verified to run.

A couple of things to decide on first render: the PowerPoint will be plainer than the RevealJS (Quarto's pptx stacks the two-column layouts); no CU branding yet.

Derek (user)

[Screenshot: quarto::quarto_render("slides/ssm_talk.qmd") -> ERROR: No valid input files passed to render.] Help me debug.

Claude (assistant)

That error almost always means Quarto can't find the file at the path given -- a **working-directory mismatch**. The file is fine (confirmed present at the project root). Check `getwd()`; it's probably not the project root.

Quickest fix -- absolute path:

```
quarto::quarto_render("/Users/briggsd/Library/CloudStorage/Dropbox/Conferences/SSM 2026/slides/ssm_talk.qmd")
```

Or `setwd()` to the project root first, or simply open the .qmd in RStudio and click Render (it sets the wd to the file's location). With no `output_format` argument it builds both HTML and pptx.

Derek (user)

That worked, thanks.

Claude (assistant)

Both `ssm_talk.html` and `ssm_talk.pptx` should now be in `slides/`. Open items whenever you're ready: eyeball the two-column and equation slides; add a CU Boulder accent / institutional title slide if you want; tighten wording or add a backup item-bank slide for Q&A.

Derek (user)

Please save the transcript of our full interaction that has led up to the production of these slides as a pdf file in the SSM 2026 root directory.

Claude (assistant)

[Produced the first version of this PDF -- a reconstructed transcript of the working session -- and saved it to the project root.]

Derek (user)

I have a question about the simulation in `verification_frontier.Rmd`.

Claude (assistant)

[Confirmed ready and recapped the file as it then stood: model setup, the two domains with capability ceilings, Fisher information plus the three self-checks, Panels A--C, the item-bank section, and the caveats.]

Derek (user)

In an unverifiable domain I'm not sure it makes sense to think of "catching" LLM output that is seductive but wrong. For the joke: if the task is to write a funny story for a dinner party, what would it mean to "catch" a story that looked funny but really was not?

Claude (assistant)

Great question --- it reframes "catch" rather than defending the literal version.

- **"Catch" was a verifiable-domain word.** There it is error-detection against a cheap, objective check. For the joke there is no such test.
 - **From detection to appraisal.** In an unverifiable domain, catching a seductive-but-wrong output = correctly judging that a fluent, joke-shaped story will fall flat for the audience, and not being captured by its polish. The "wrong" is **won't-achieve-the-goal**; the seduction is that LLMs produce text that pattern-matches "funny" without being funny for the specific room.
 - **The deep point.** To **measure** whether someone caught correctly you need a scoring key --- and for the joke that key is the **deferred, social verdict** (tell it; the audience laughs or not). So "unverifiable" is really **expensively / deferredly verifiable**; the real axis is the ****cost and latency of verification****, not a binary. A genuinely unverifiable domain (pure preference, no convergent verdict) would be **unmeasurable**.
 - **Why this strengthens the thesis.** The same deferred, costly verification both caps trainable capability (so the construct persists) and is the only scoring key for measuring the human --- so critical literacy matters most where verification is expensive and is hardest to measure for the same reason. Three options: (1) relabel only, no math change; (2) relabel plus an upper asymptote $u < 1$ in deferred domains; (3) leave the math, add a caveat. Lean 1 now, 2 if you want the model to carry the point.
-

Derek (user)

Let's do 1. Please also fix some math notation that didn't render correctly in section 8 when showing the distribution of item difficulty.

Claude (assistant)

Applied option 1 across the Rmd and the slides. "Catch" now reads as **over-riding a plausible-but-inadequate answer** --- error-detection where verification is cheap, appraisal where it is deferred. The "Two domains" section is reframed from a verifiable/unverifiable binary to the ****cost and latency of verification****, with the deferred verdict doing double duty (it caps capability ***and*** supplies the scoring key) and the corollary that a domain with no convergent verdict is unmeasurable, not merely costly. No math changed. Section 8 fix: the item-difficulty line that read $N(\delta_{of_s}(s), \sigma)$ --- whose escaped underscores inside a text-mode block were the culprit --- now reads $N(\delta(s), \sigma)$, which is cleaner and the correct symbol anyway. Verified the Rmd still knits and no fragile math commands remain.

Derek (user)

Is it correct to say that the plot in Panel B is specific to a cheaply verifiable task, or does the distinction not matter here? In Panel A we contrast $v = 1$ vs $v = 0$, but in Panel B that contrast is lost.

Claude (assistant)

Right to interrogate it. Panel B's math depends on c and s only --- v is never an argument --- so the curves are domain-agnostic; in that narrow sense the distinction "doesn't matter." But v controls ***which*** c a domain can reach, and $c = 0.90$ is attainable only by a cheaply-verifiable task (unverifiable caps at ~ 0.55). So the "frontier" curve implicitly told the verifiable story and the unverifiable frontier was missing --- the contrast was hidden, and the omission matters because the unverifiable frontier ($c = 0.55$) lies almost on top of the early-model curve. Fix: show three curves at the capabilities the two roads actually occupy --- early (0.50), unverifiable frontier (0.55), verifiable frontier (0.96) --- restoring the same v -contrast as Panel A, sliced along s .

Derek (user)

Yes [make that edit].

Claude (assistant)

Rebuilt Panel B with the three curves in the Rmd, the slides, and the standalone Python PNG. The unverifiable frontier (red, $c = 0.55$) sits just under the early model (grey); only the verifiable frontier (blue, $c = 0.96$) flattens toward zero. Verified all render.

Derek (user)

The text for the new plot in Panel B is running into the curves. Get rid of the text and just use $c = .50$, $c = .55$, $c = .96$. I can explain how this relates to v in my talk.

Claude (assistant)

Shortened the Panel B legend to just $c = 0.50$, $c = 0.55$, $c = 0.96$ (Rmd + Python), which also keeps it tucked in the corner clear of the peak. The colours still encode the domains for your own reference.

Derek (user)

After introducing the measurement model in section 2, include a plot of a hypothetical item characteristic curve. This will make it easier to explain how to interpret the c and $\Delta(s)$ parameters to my audience.

Claude (assistant)

Added a subsection "Reading the parameters: an item characteristic curve" right after the model definition. The ICC shows c as the **floor** (success with zero critical literacy, rising with model capability), $\Delta(s)$ as the **location** (the literacy level for a coin-flip over-ride), and a second, more-seductive curve shifted right. Annotations sit in empty regions --- I rendered it and checked for the run-into-curves problem you flagged on Panel B. The Rmd renders cleanly.

Derek (user)

Yes [add it to the slides].

Claude (assistant)

Added a matching ICC slide ("Reading the parameters") after "A measurement model" in the deck, using the deck's domain colours, with a one-line caption reinforcing both parameters. All slide chunks verified.

Derek (user)

Can you post the Rmd and HTML rendering (verification_frontier.Rmd and verification_frontier.html) to a new folder labelled "SSM 2026 Presentation" in my GitHub repo (Measurement-and-Psychometrics)?

Claude (assistant)

Published both to a new "SSM 2026 Presentation" folder on master. Along the way a `grep -c` quirk made the HTML look stale (it is unreliable on the 2.7 MB near-single-line HTML); a proper visible-text extraction confirmed both the Rmd and the HTML were fully current, so nothing stale went up. `gh` was not installed, so I pushed over SSH under your identity, directly to master --- matching your request and the repo's existing spaced-folder convention. Verified the files landed and cleaned up the temporary clone.

Derek (user)

Let's update the transcript with today's work and post it to the repository. Don't post the slides --- I'm creating "Briggs_SSM_Measurement_AI" by hand and am not ready to post those yet.

Claude (assistant)

[Updated this transcript with today's session, regenerated the PDF, and posted it to the "SSM 2026 Presentation" folder in the repo. Did not post the slide deck, per your request.]
